

Bank Marketing Effectiveness Prediction

using various Machine Learning Models

Project submitted to the
SRM University – AP, Andhra Pradesh
for the partial fulfillment of the requirements to award the degree of

Bachelor of Technology

In

Computer Science and Engineering
School of Engineering and Sciences

Submitted by

Kothuri Guru Sai Keerthana		AP21110011578
Lakku Manohar Reddy		AP21110011581
Amancha Phalgun		AP21110011587
Botla Venkata Ramamohan Rao		AP21110011598



Under the Guidance of
Dr. Satya Pramod Jammy
SRM University-AP
Neerukonda, Mangalagiri, Guntur
Andhra Pradesh – 522 240
May, 2024




Table of Contents

Abstract.....	Error! Bookmark not defined.
Abbreviations	Error! Bookmark not defined.
List of Figures	Error! Bookmark not defined.
List of Equations	Error! Bookmark not defined.
1. Introduction	Error! Bookmark not defined.
2.Dataset Description	Error! Bookmark not defined.
3. Machine Learning Models	2Error! Bookmark not defined.
3.1. KNN (K-Nearest Neighbours)	2Error! Bookmark not defined.
3.2 Logistic Regression	24
3.3 SVM (Support Vector Machine)	24
3.4 K means	25
Results.....	26
4.1 Knn.....	26
4.2. Logistic Regression.....	27
4.3 SVM	28
4.4 Kmeans	29
5. Concluding Remarks	31
References	35

Abstract

This project explores the application of machine learning techniques to predict the effectiveness of bank marketing campaigns using a dataset from a Portuguese banking institution. With a focus on classification, the goal is to develop models that can accurately classify clients' responses to campaigns as positive or negative. The dataset contains various input variables such as age, job, marital status, education, and financial indicators.

Initial data analysis involved computing descriptive statistics and visualizations to understand the relationships between variables. Outliers were addressed using interquartile range, and missing values were imputed or features were eliminated if they contained over 50% null values. Insights from the analyses revealed patterns such as age group preferences, job categories, marital status, education levels, and loan statuses affecting subscription to term deposits.

The implemented algorithms, k-Means clustering, SVM, KNN, and Logistic Regression, underwent training and evaluation. Cross-validation techniques were employed to enhance model performance.

Our goal is to develop robust ML models that are accurate. To evaluate the performance of the models, we employ metrics such as accuracy, precision, recall, and F1-score.

The findings offer valuable insights into campaign effectiveness and demonstrate the potential of machine learning in optimizing marketing strategies. Though only a subset of algorithms was implemented, the study lays the groundwork for future analyses leveraging additional techniques to further enhance prediction accuracy.

Abbreviations

EDA	Exploratory Data Analysis
TP	True Positive
TN	True Negative
FP	False Positive
FN	False Negative
KNN	K Nearest Neighbours
SVM	Support Vector Machine

List of Figures

Figure 1	Confusion Matrix	v
Figure 2	KNN Result	5
Figure 3	Logistic Regression Result	6
Figure 4	Support Vector Machine Result.....	7

List of Equations

Confusion Matrix:

		PREDICTED	
		Positive	Negative
ACTUAL	Positive	TRUE POSITIVE	FALSE NEGATIVE
	Negative	FALSE POSITIVE	TRUE NEGATIVE

Figure 1 Confusion Matrix

$$\text{Accuracy} = \frac{TP+TN}{TP+TN+FP+FN}$$

$$\text{Precision} = \frac{TP}{TP+FP}$$

$$\text{Recall} = \frac{TP}{TP+FN}$$

$$\text{F1 Score} = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

KNN Distance Metrics:

1. Euclidean Distance :

$$\text{distance}(x, X_i) = \sqrt{\sum_{j=1}^d (x_j - X_{ij})^2}$$

2. Manhattan Distance:

$$d(x, y) = \sum_{i=1}^n |x_i - y_i|$$

3. Minkowski Distance:

$$d(x, y) = (\sum_{i=1}^n (x_i - y_i)^p)^{\frac{1}{p}}$$

Logistic Regression:

1.Sigmoid Function: $\sigma = \frac{1}{(1+e^{-x})}$

2.Cost Function: $cost = -\frac{1}{m} \sum_{i=1}^m [y * \log(a) + (1 - y) * \log(1 - a)]$

3.Gradient Descent:

$$dW = \frac{\partial COST}{\partial W} = (A - Y) * X^T \text{ shape (1 x n)}$$

$$dB = \frac{\partial COST}{\partial B} = (A - Y)$$

$$W = W - \alpha * dW^T$$

$$B = B - \alpha * dB$$

1. Introduction

Bank marketing campaigns play a crucial role in financial institutions' efforts to attract and retain customers. Understanding the factors that influence the effectiveness of these campaigns is essential for optimizing marketing strategies and maximizing returns on investment. In this context, machine learning techniques offer a promising avenue for predicting client responses to marketing initiatives. Leveraging a dataset provided by a Portuguese banking institution, this project delves into the predictive modeling of bank marketing campaign effectiveness. By analyzing a wide range of client demographics and financial indicators, the aim is to develop accurate classification models capable of discerning whether clients are likely to subscribe to term deposits based on campaign outreach.

The dataset comprises a rich array of input variables, including age, job type, marital status, education level, and financial status indicators such as balance and loan status. Initial exploratory data analysis reveals intriguing insights into the relationships between these variables and clients' propensity to subscribe to term deposits. From age group preferences to the influence of job categories and loan statuses, the data illuminates various factors that may impact campaign effectiveness. Such insights serve as a foundation for building robust predictive models that can inform targeted marketing strategies tailored to specific client demographics and financial profiles.

Addressing data preprocessing challenges, including missing values and outliers, is paramount to ensure the integrity and efficacy of the predictive modeling process. Through techniques such as imputation, feature elimination, and outlier treatment using the interquartile range, the dataset is refined for further analysis. Moreover, the presence of class imbalance, with a significantly higher number of clients not subscribing to term deposits compared to those who do, necessitates the adoption of oversampling techniques like Synthetic Minority Oversampling Technique (SMOTE) to mitigate bias and enhance model performance.

With a focus on implementing key machine learning algorithms, namely kMeans clustering, Support Vector Machine (SVM), K Nearest Neighbors (KNN), and Logistic Regression, this project aims to provide a comprehensive analysis of bank marketing campaign effectiveness prediction. By leveraging the strengths

of these algorithms and employing cross-validation techniques to refine model performance, the study endeavors to offer actionable insights for financial institutions seeking to optimize their marketing strategies and drive better campaign outcomes.


2.Dataset Description

The dataset utilized in this project is obtained from a Portuguese banking institution and contains crucial information related to bank marketing campaigns. It comprises 45211 observations and 17 columns, each providing valuable insights into various aspects of client interactions and campaign outcomes. The dataset includes demographic details such as age, job type, marital status, and education level, along with financial indicators like account balance, loan status, and contact method. Additionally, temporal variables such as the day and month of contact, as well as campaign-specific metrics like call duration and the number of contacts, are incorporated into the dataset.


The data set contained details about bank marketing campaigns. Descriptive statistics were computed for each variable as part of the analysis, and visualizations were made to investigate the relationships between the various variables. We created a number of graphs, such as a distplot, count plot, bar plot, pair plot to gain insight from the dataset.

Categorical variables within the dataset, including job, marital status, education, contact method, and outcome of the previous marketing campaign, offer valuable insights into clients' socio-demographic backgrounds and previous interactions with the bank. These variables provide context for understanding client behavior and preferences, which are essential for predicting campaign effectiveness. Furthermore, numerical variables such as age, account balance, and call duration offer quantitative measures of clients' financial status and engagement with the marketing campaign. The target variable, denoted as 'y', indicates whether a client subscribed to a term deposit following the marketing campaign, facilitating the classification task of predicting campaign effectiveness.


Data Preprocessing and handling Outliers:




Histogram of Age:




Boxplot of marital




Histogram of marital

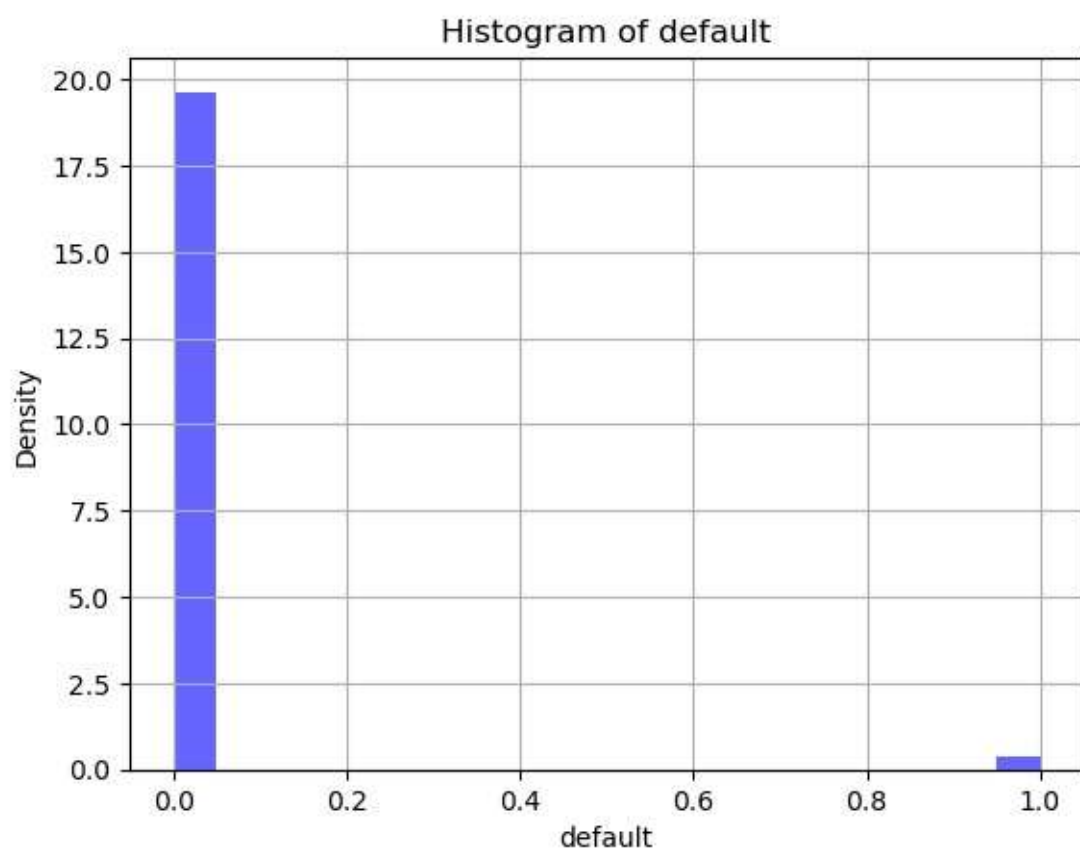
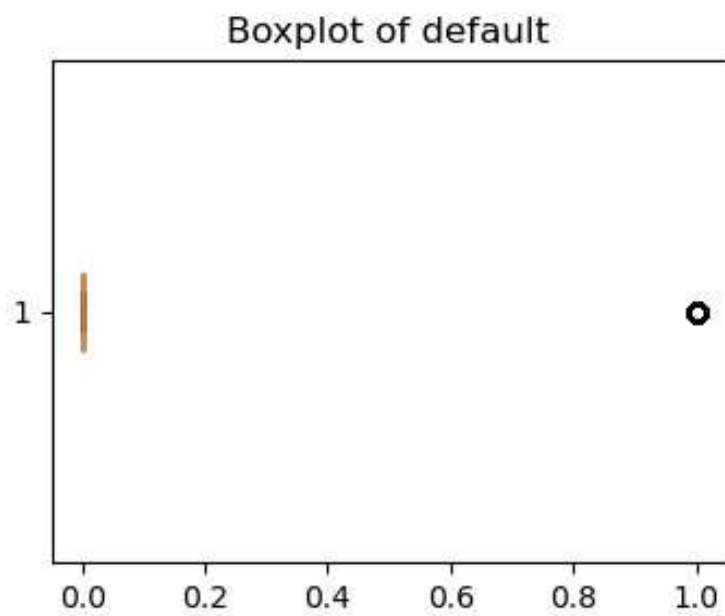


Boxplot of education




Histogram of education




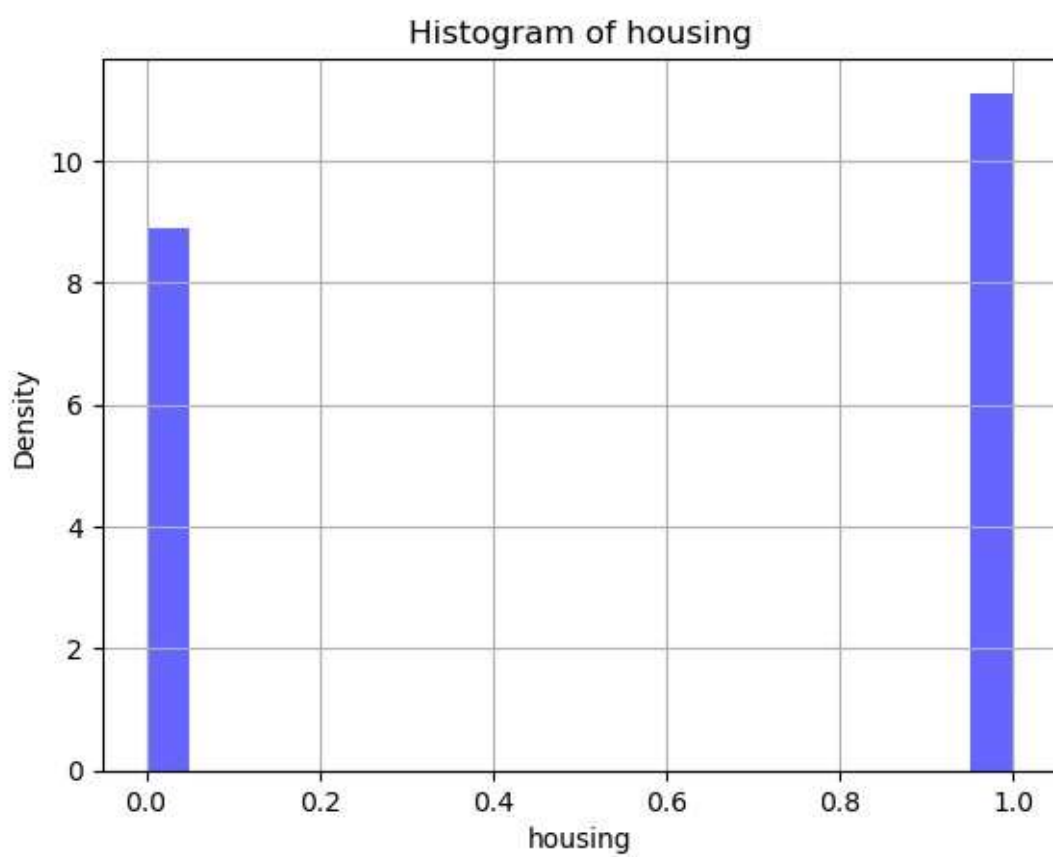
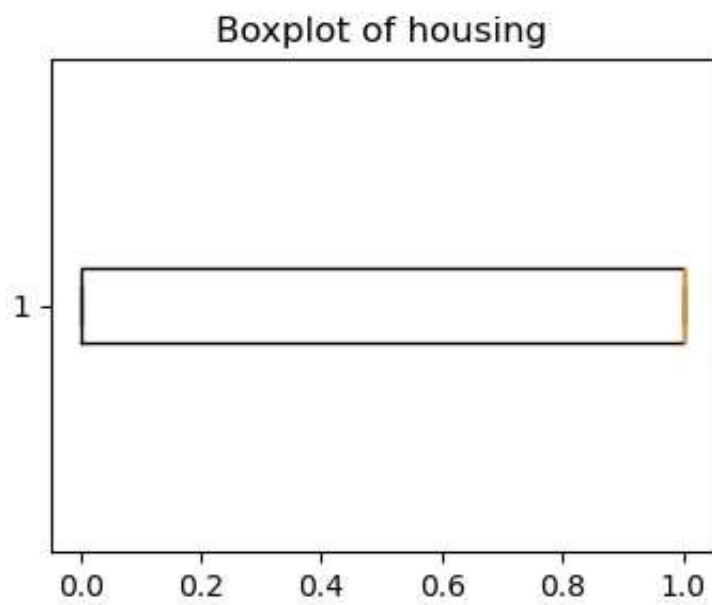


Boxplot of balance




Histogram of balance




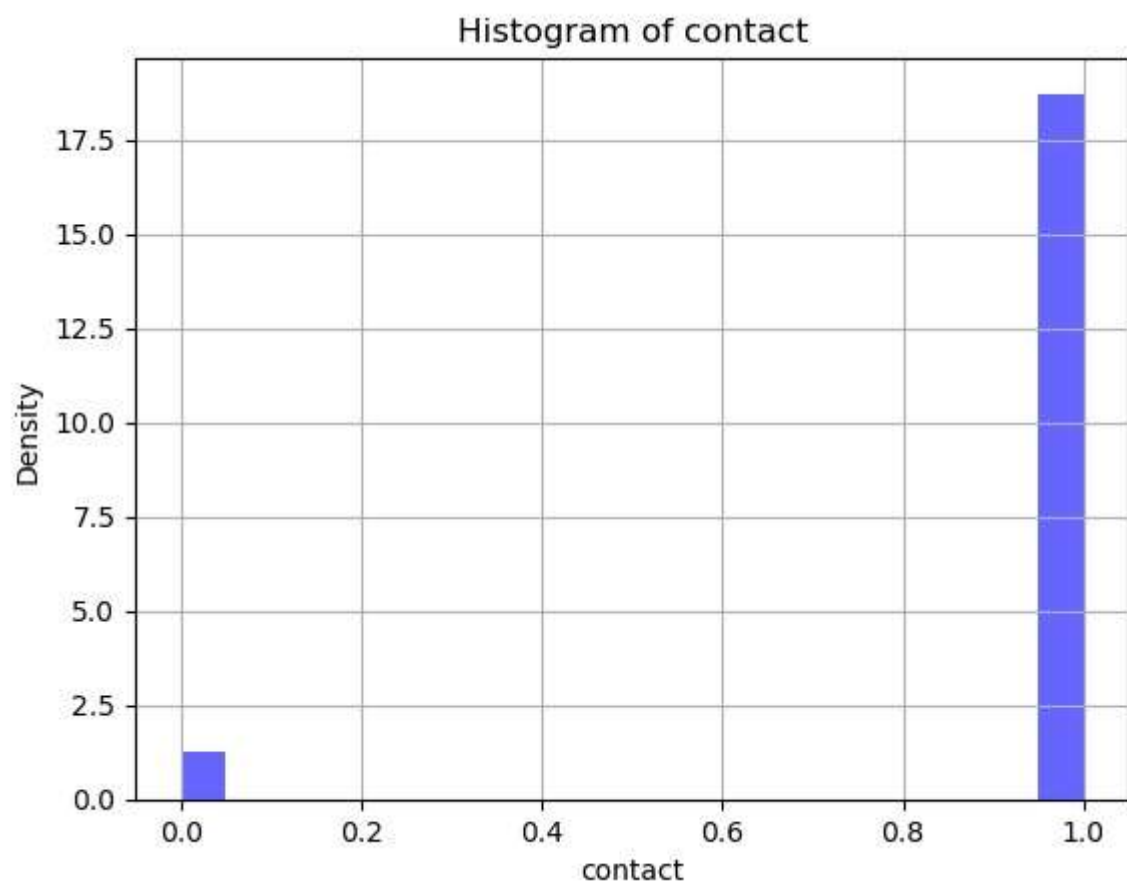
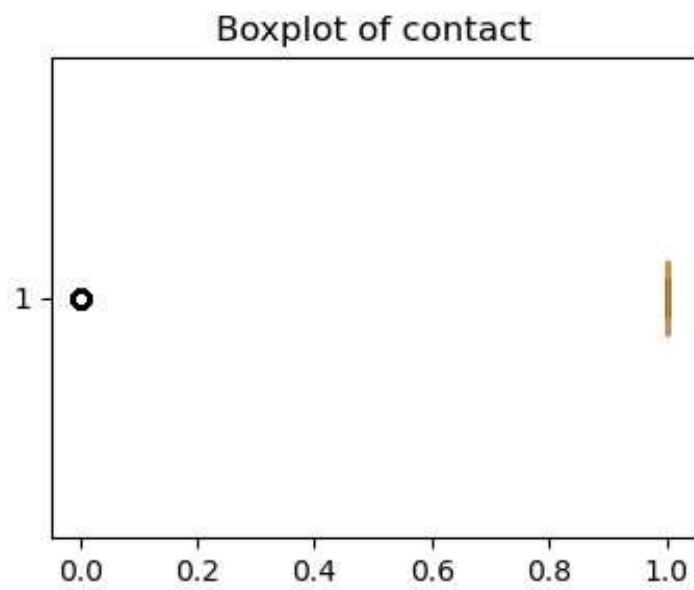


Boxplot of loan




Histogram of loan







Boxplot of day




Histogram of day




Boxplot of duration



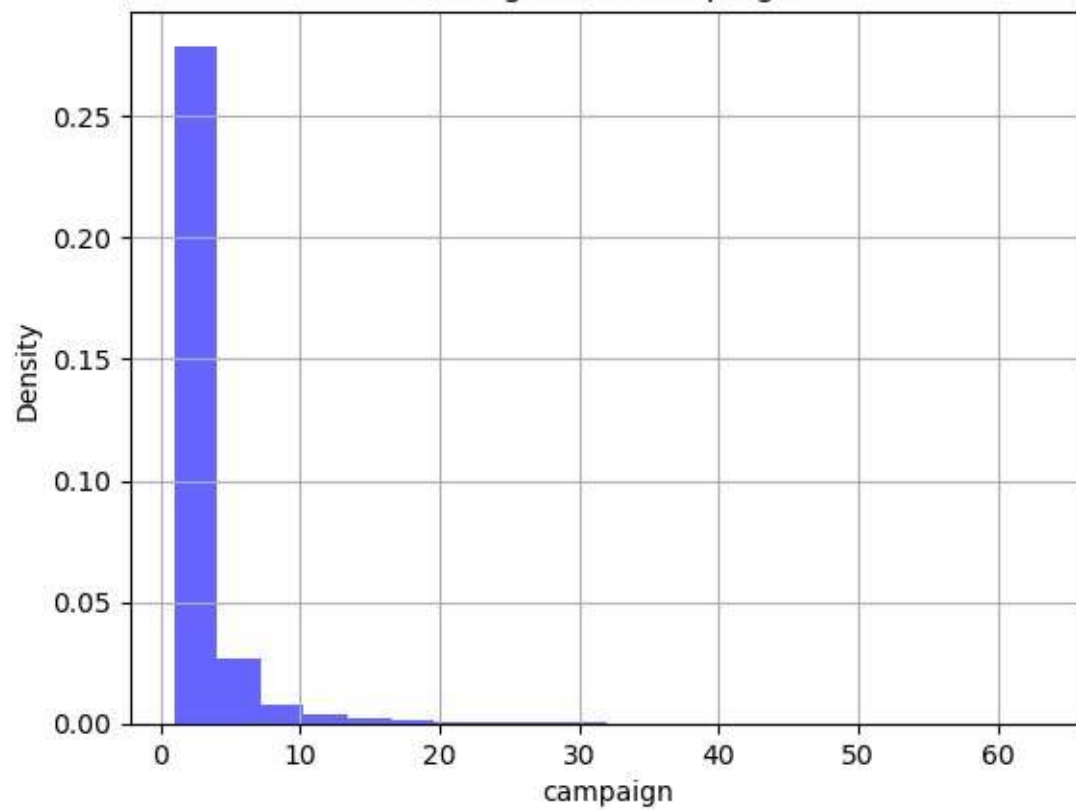
Histogram of duration

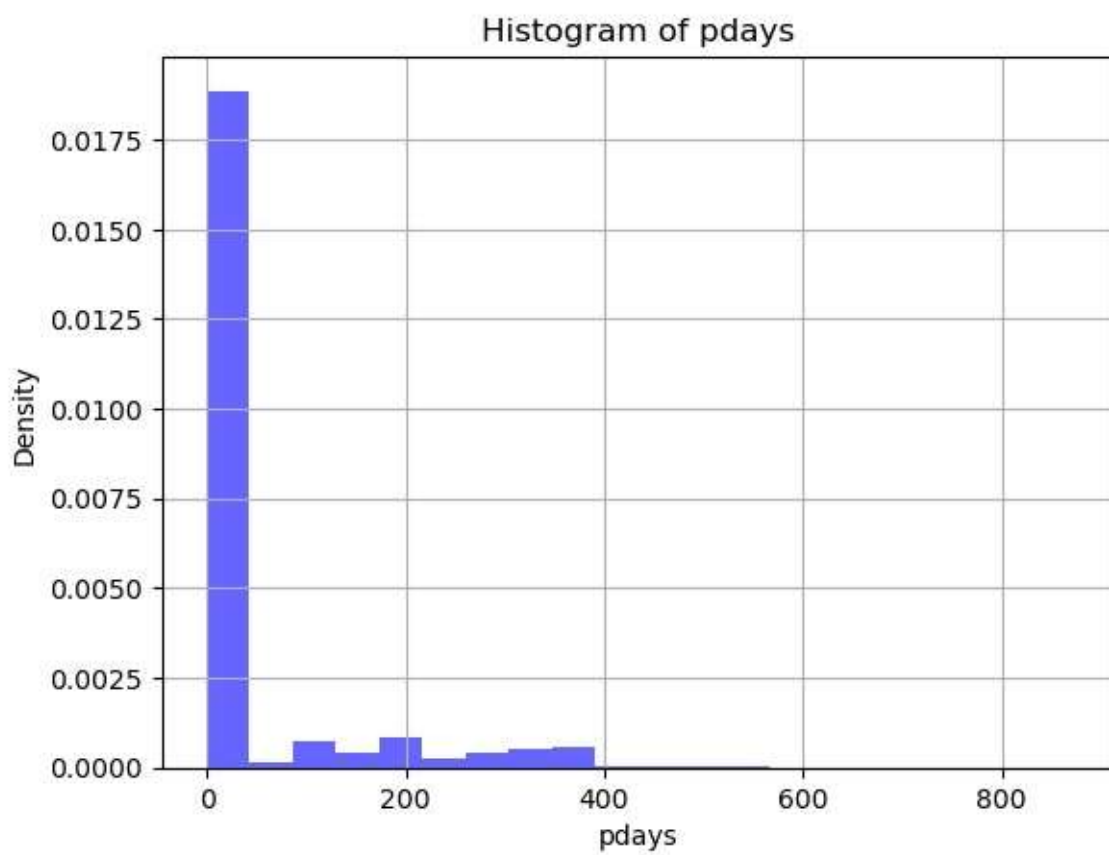
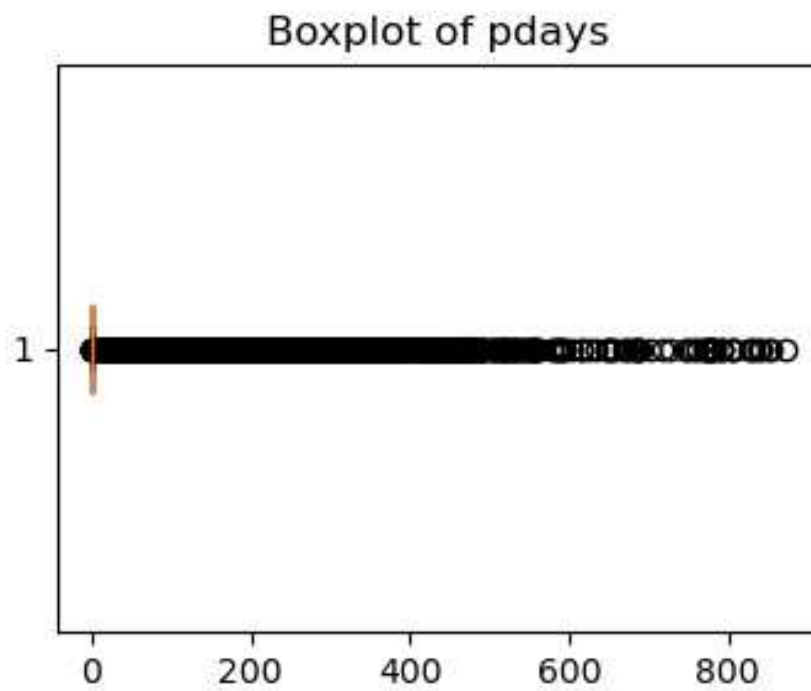


Boxplot of campaign




Histogram of campaign




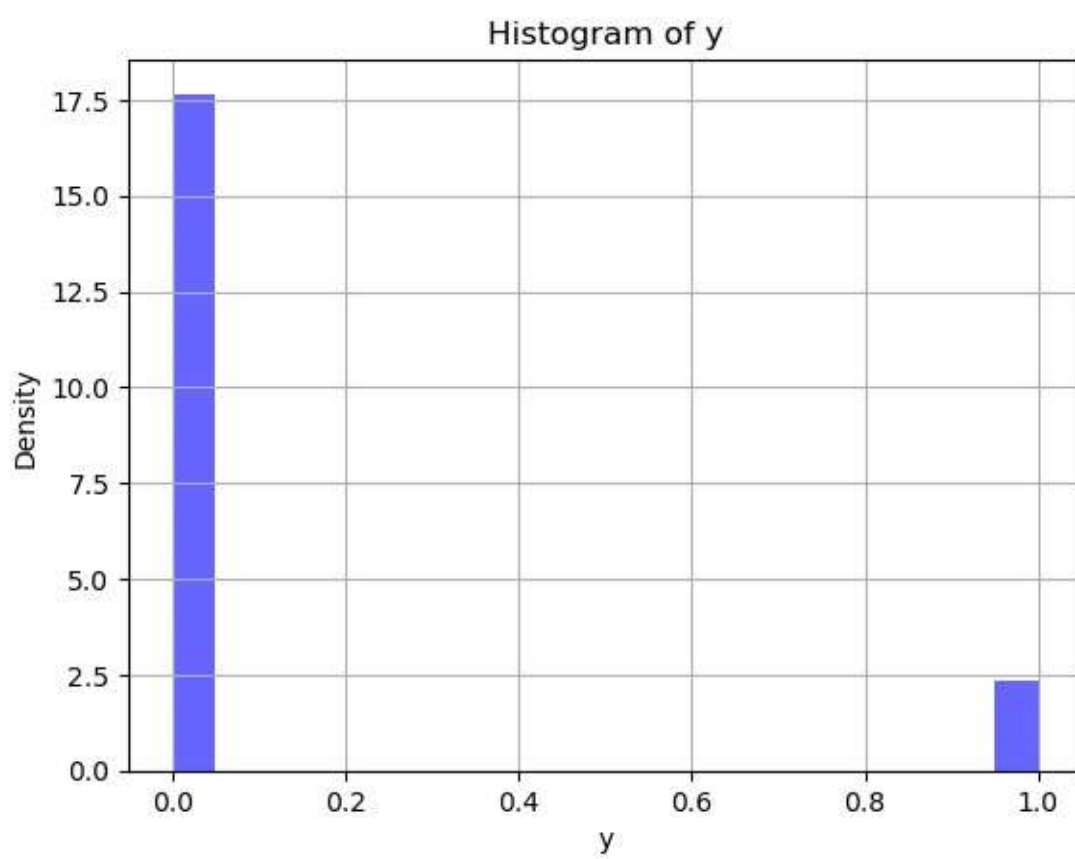
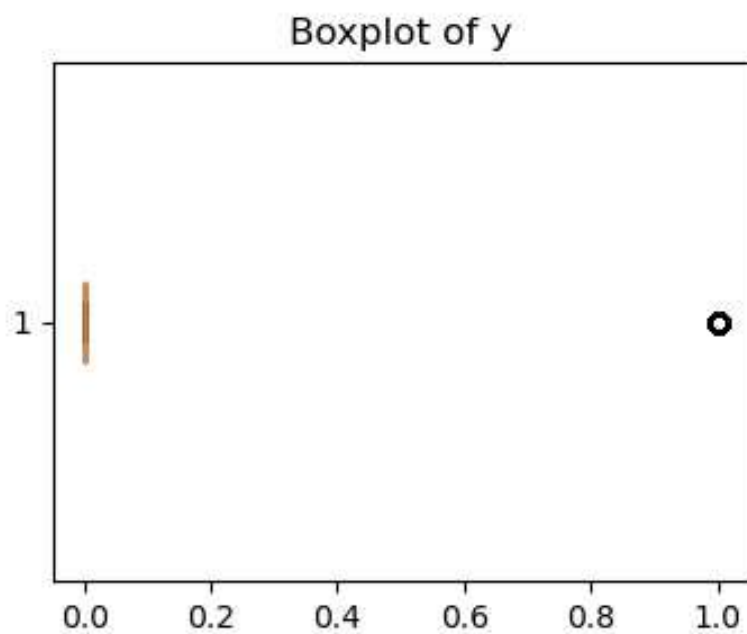


Boxplot of previous



Histogram of previous





3. Machine Learning Models

3.1. KNN (K-Nearest Neighbours)

In the domain of predicting the effectiveness of bank marketing campaigns, the K-Nearest Neighbors (KNN) algorithm emerges as a promising methodology. Utilizing a feature space defined by diverse attributes like demographics, financial behaviors, and historical engagement patterns, KNN endeavors to classify whether a client is likely to subscribe to a term deposit or not. This process involves comparing the feature vectors of new clients with those of existing ones in the training dataset, identifying the K nearest neighbors based on a specified distance metric, and assigning the majority class label among these neighbors to the new client. Despite its simplicity, KNN can yield valuable insights into the dynamics of campaign effectiveness, particularly in scenarios characterized by nonlinear feature-churn relationships or where interpretability is crucial. However, it's imperative for practitioners to fine-tune parameters such as the value of K and carefully consider computational demands, particularly in scenarios involving large-scale datasets. Nevertheless, KNN remains a versatile tool in the arsenal of predictive analytics, providing an intuitive and interpretable approach to predicting the effectiveness of bank marketing campaigns.

3.2 Logistic Regression

Logistic Regression is a cornerstone of our project for forecasting the efficacy of bank marketing campaigns. In this context, it functions by estimating the probability of a client subscribing to a term deposit based on a multitude of predictor variables. These variables encompass demographic details, financial indicators, and temporal factors, collectively providing insight into client behavior and preferences.

One of the primary advantages of logistic regression in our project lies in its simplicity and interpretability. By modeling the relationship between features and campaign effectiveness, logistic regression offers a clear understanding of the factors influencing client decisions to subscribe to term deposits. This transparency is invaluable for banks aiming to tailor their marketing strategies effectively.

Furthermore, logistic regression's ability to handle nonlinear relationships between features and campaign outcomes makes it well-suited for our predictive task. By analyzing the coefficients associated with each feature, we can discern their relative importance in influencing client behavior. This insight aids in prioritizing marketing efforts and optimizing resource allocation for maximum impact.

3.3 SVM (Support Vector Machine)

In our project, Support Vector Machine (SVM) emerges as a potent tool for forecasting the effectiveness of bank marketing campaigns. SVM operates by delineating the optimal hyperplane that segregates the data into distinct classes, in this case, clients who subscribe to term deposits and those who do not. By maximizing the margin between these classes, SVM seeks to identify the most effective decision boundary based on various client features.

In the realm of campaign effectiveness prediction, SVM endeavors to discern the optimal decision boundary that segregates clients likely to subscribe to term deposits from those less inclined to do so. This is achieved by leveraging various client features such as demographics, financial behaviors, and historical engagement patterns. By mapping the data into a higher-dimensional space using kernel functions, SVM can effectively capture complex relationships between features and campaign outcomes, facilitating the identification of nonlinear decision boundaries.

The flexibility of SVM to capture intricate patterns in client data makes it particularly suitable for our predictive task. By accurately delineating the decision boundary, SVM aims to generalize well to unseen data, enabling precise predictions regarding client subscription behavior. However, it's essential to note that SVMs can be computationally intensive, especially with large datasets, and require meticulous parameter tuning for optimal performance. Nonetheless, with careful parameter optimization and feature engineering, SVMs hold the potential to offer valuable insights into campaign effectiveness, aiding financial institutions in optimizing their marketing strategies and maximizing subscription rates.

3.4 Kmeans

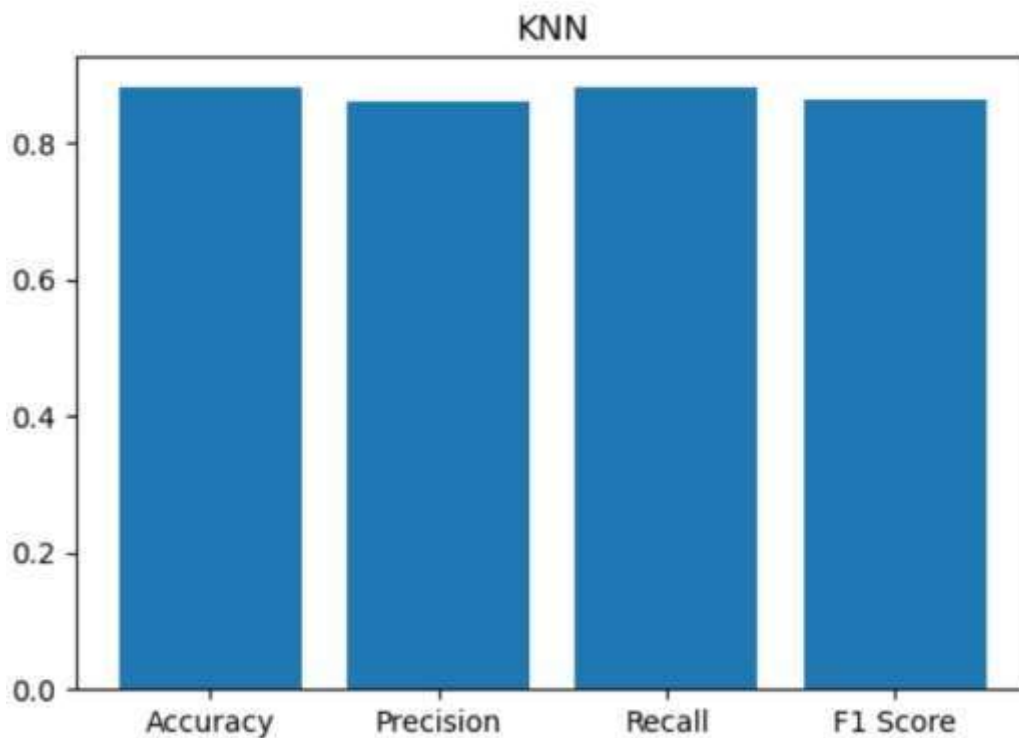
In the context of campaign effectiveness prediction, K-Means seeks to identify distinct groups of clients with similar characteristics and behaviors. By clustering clients based on features such as demographics, financial indicators, and past interaction patterns, K-Means provides valuable insights into heterogeneous

client segments that may exhibit varying propensities to subscribe to term deposits.

The flexibility of K-Means to uncover hidden patterns in client data makes it particularly suitable for our predictive task. By grouping clients into clusters with similar subscription behaviors, K-Means facilitates the identification of key client segments that are more likely to respond positively to marketing campaigns. However, it's essential to note that K-Means clustering requires careful consideration of the number of clusters (K) and may not perform optimally in the presence of non-linear relationships between features.

Results:

4.1. KNN




Confusion Matrix:

```
[[7668  276]
 [ 803  296]]
```

	precision	recall	f1-score	support
0	0.91	0.97	0.93	7944
1	0.52	0.27	0.35	1099
accuracy			0.88	9043
macro avg	0.71	0.62	0.64	9043
weighted avg	0.86	0.88	0.86	9043

4.2. Logistic Regression




Confusion Matrix::

```
[[7834  110]
 [ 897  202]]
```

	precision	recall	f1-score	support
0	0.90	0.99	0.94	7944
1	0.65	0.18	0.29	1099
accuracy			0.89	9043
macro avg	0.77	0.58	0.61	9043
weighted avg	0.87	0.89	0.86	9043

Figure 3 Logistic Regression Result

4.3 SVM




Confusion Matrix :

```
[[7941  3]
 [1097  2]]
```

	precision	recall	f1-score	support
0	0.88	1.00	0.94	7944
1	0.40	0.00	0.00	1099
accuracy			0.88	9043
macro avg	0.64	0.50	0.47	9043
weighted avg	0.82	0.88	0.82	9043

Figure 4 Support Vector Machine Result

4.4 K-Means:

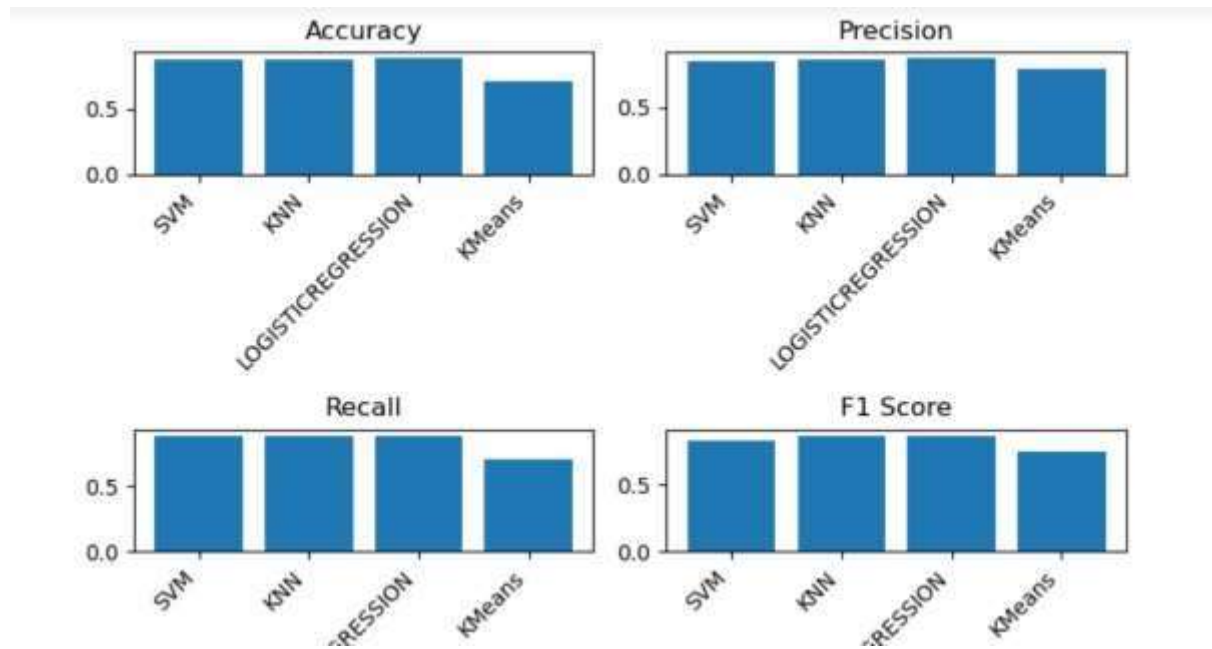


Confusion Matrix:

```
[[7944  0]
 [1099  0]]
```

Classification report:

	precision	recall	f1-score	support
0	0.88	1.00	0.94	7944
1	0.00	0.00	0.00	1099
accuracy			0.88	9043
macro avg	0.44	0.50	0.47	9043
weighted avg	0.77	0.88	0.82	9043



Concluding Remarks

Model	Accuracy	Precision	Recall	F1 Score
SVM	0.879575	0.850356	0.879575	0.827157
KNN	0.881455	0.859081	0.881455	0.864535
LOGISTICREGRESSION	0.890191	0.869367	0.890191	0.867404
KMeans	0.711158	0.790686	0.711158	0.745773

The analysis provides valuable insights into the factors influencing customer subscription to bank term deposits and highlights crucial trends in customer behavior. Here are some conclusion remarks based on the observations:

1. Demographic Patterns: Understanding the age distribution and occupation types of clients reveals nuances in subscription behavior. The analysis provides valuable insights into the factors influencing customer subscription to bank term deposits and highlights crucial trends in customer behavior. Here are some conclusion remarks based on the observations:

1. Demographic Patterns: Understanding the age distribution and occupation types of clients reveals nuances in subscription behavior. While the average client age spans a wide range, the peak subscription age range suggests a specific demographic preference for term deposits, possibly influenced by financial stability and life stage considerations.

2. Marital Status and Education: Marital status and education levels play significant roles in subscription likelihood, indicating potential correlations between financial literacy, stability, and commitment levels.

3. Financial Obligations: The absence of defaults and housing loans positively influences subscription rates, while the presence of personal loans or multiple loan types deters subscription, reflecting the impact of financial constraints and risk aversion.

4. Communication Channels and Frequency: Contact through cellular phones appears to be more effective in prompting subscriptions, with a diminishing return observed beyond three contact attempts. This underscores the importance of targeted and timely communication strategies.

5. Seasonal Variations: Seasonality affects subscription rates, with May standing out as a particularly high-performing month. Understanding these seasonal fluctuations can inform marketing strategies and resource allocation.
6. Occupational and Educational Profiles: Clients with managerial roles and higher education levels exhibit higher subscription rates, suggesting a link between financial sophistication and propensity to invest in term deposits.
7. Engagement Metrics: Longer phone conversations correlate with higher subscription rates, indicating a potential relationship between engagement and interest in financial products.
8. Debt Status and Financial Freedom: Debt-free clients are more likely to subscribe, emphasizing the importance of financial stability and capacity for investment.
9. Model Performance and Recommendations: Machine learning models, particularly Logistic Regression and SVM, show promising accuracy rates in predicting churn. Leveraging ensemble methods and incorporating additional features could further enhance predictive power and robustness.
10. Business Implications: Businesses can use these insights to refine customer retention strategies, tailor communication channels, and optimize resource allocation. Proactively addressing churn risks and fostering stronger customer relationships are essential for long-term success in a competitive market landscape.

In conclusion, by leveraging the multifaceted insights gleaned from this analysis, businesses can make informed decisions to mitigate churn, enhance customer satisfaction, and drive sustainable growth. Continuous monitoring, adaptation, and innovation in predictive analytics remain critical for staying ahead in an ever-evolving market environment. While the average client age spans a wide range, the peak subscription age range suggests a specific demographic preference for term deposits, possibly influenced by financial stability and life stage considerations.

2. Marital Status and Education: Marital status and education levels play significant roles in subscription likelihood, indicating potential correlations between financial literacy, stability, and commitment levels.
3. Financial Obligations: The absence of defaults and housing loans positively influences subscription rates, while the presence of personal loans or multiple loan types deters subscription, reflecting the impact of financial constraints and risk aversion.
4. Communication Channels and Frequency: Contact through cellular phones appears to be more effective in prompting subscriptions, with a

diminishing return observed beyond three contact attempts. This underscores the importance of targeted and timely communication strategies.

5. **Seasonal Variations:** Seasonality affects subscription rates, with May standing out as a particularly high-performing month. Understanding these seasonal fluctuations can inform marketing strategies and resource allocation.
6. **Occupational and Educational Profiles:** Clients with managerial roles and higher education levels exhibit higher subscription rates, suggesting a link between financial sophistication and propensity to invest in term deposits.
7. **Engagement Metrics:** Longer phone conversations correlate with higher subscription rates, indicating a potential relationship between engagement and interest in financial products.
8. **Debt Status and Financial Freedom:** Debt-free clients are more likely to subscribe, emphasizing the importance of financial stability and capacity for investment.
9. **Model Performance and Recommendations:** Machine learning models, particularly Logistic Regression and SVM, show promising accuracy rates in predicting churn. Leveraging ensemble methods and incorporating additional features could further enhance predictive power and robustness.
10. **Business Implications:** Businesses can use these insights to refine customer retention strategies, tailor communication channels, and optimize resource allocation. Proactively addressing churn risks and fostering stronger customer relationships are essential for long-term success in a competitive market landscape.

In conclusion, by leveraging the multifaceted insights gleaned from this analysis, businesses can make informed decisions to mitigate churn, enhance customer satisfaction, and drive sustainable growth. Continuous monitoring, adaptation, and innovation in predictive analytics remain critical for staying ahead in an ever-evolving market environment.

References :

1. Han, J., Kamber, M.: Data Mining: Concepts and Techniques. Morgan Kaufmann, Burlington (2000).
2. Pujari, A.K.: Data Mining Techniques, 1st edn. Universities Press (India) Private Limited, Hyderabad (2001)

ml-project

May 9, 2024

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
[2]: df=pd.read_csv('bank-full.csv',sep=";")
df
```

```
[2]:      age      job  marital  education  default  balance  housing  loan  \
0      58  management  married    tertiary    no      2143    yes    no
1      44  technician  single  secondaryno    29      yes    no 2  33
      entrepreneur  married  secondary    no      2      yes  yes 3
      47  blue-collar    married    unknown    no      no      1506
      yes    no
4      33    unknown    single    unknown    no      1      no    no
...  ...
45206  51  technician  married    tertiary    no      825    no    no
45207  71  retired  divorcedprimary    no    1729    no    no
45208  72  retired  married  secondary    no    5715    no    no
45209  57  blue-collar    married  secondary    no      668    no    no
45210  37  entrepreneur  married  secondary    no    2971    no    no
      contact  day  month  duration  campaign  pdays  previous  poutcome  y
0      unknown    5    may      261          1     -1          0  unknown    no
1      unknown    5    may      151          1     -1          0  unknown    no
2      unknown    5    may       76          1     -1          0  unknown    no
3      unknown    5    may       92          1     -1          0  unknown    no
4      unknown    5    may      198          1     -1          0  unknown    no
...  ...
45206  cellular   17    nov      977          3     -1          0  unknown  yes
45207  cellular   17    nov      456          2     -1          0  unknown  yes
45208  cellular   17    nov     1127          5    184          3  success  yes
45209  telephone  17    nov      508          4     -1          0  unknown    no
45210  cellular   17    nov      361          2    188         11    other    no
[45211 rows x 17 columns]
```

```
[3]: df.columns
[3]: Index(['age', 'job', 'marital', 'education', 'default', 'balance',
'housing',
        'loan', 'contact', 'day', 'month', 'duration', 'campaign',
        'pdays',
        'previous', 'poutcome', 'y'],
        dtype='object')
[4]: df.info()
```

```
<class
'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to
45210 Data columns (total 17
columns):
#   Column      Non-Null Count
   Dtype
---  -
0   age         45211 non-null
   int64
1   job         45211 non-null
   object
2   marital     45211 non-null
   object
3   education   45211 non-null
   object
4   default     45211 non-null
   object
5   balance     45211 non-null
   int64
6   housing     45211 non-null
   object
7   loan        45211 non-null
   object
8   contact     45211 non-null
   object
9   day         45211 non-null
   int64
10  month       45211 non-null
   object
11  duration    45211 non-null
   int64
12  campaign    45211 non-null
   int64
13  pdays       45211 non-null
   int64
```



```

14 previous    45211 non-null
               int64
15 poutcome    45211      non-null
               object
16 y           45211      non-null
               object
dtypes: int64(7), object(10)
memory usage: 5.9+ MB

```

```
[5]: df.head()
```

```

[5]:   age      job marital education default balance housing loan \
0     58  management married   tertiary    no    2143   yes
no 1  44  technician single secondaryno    29    yes   no
2    33 entrepreneur married secondary    no    2    yes yes
3    47 blue-collar married   unknown    no   1506   yes   no
4    33 unknown single   unknown    no    1    no    no

      contact day month duration campaign pdays previous poutcome y
0 unknown    5   may    261         1    -1         0 unknown no
1 unknown    5   may    151         1    -1         0 unknown no
2 unknown    5   may    76         1    -1         0 unknown no
3 unknown    5   may    92         1    -1         0 unknown no
4 unknown    5   may   198         1    -1         0 unknown no

```

```
[6]: df.tail()
```

```

[6]:   age      job marital education default balance housing loan \
45206  51  technician married   tertiary    no    825   no   no
45207  71  retired divorcedprimary    no   1729   no   no
45208  72  retired married secondary    no   5715   no   no
45209  57  blue-collar married secondary    no    668   no   no
45210  37 entrepreneur married secondary    no   2971   no   no

      contact day month duration campaign pdays previous poutcome y
45206 cellular  17   nov    977         3    -1         0 unknown yes
45207 cellular  17   nov    456         2    -1         0 unknown yes
45208 cellular  17   nov   1127         5   184         3 success yes
45209 telephone 17   nov    508         4    -1         0 unknown no
45210 cellular  17   nov    361         2   188        11 other no

```

```
[7]: df.shape
```

```
[7]: (45211, 17)
```

```
[8]: df.describe()
```

```

[8]: age      balance      day      duration      campaign \ count 45211.000000
      45211.000000 45211.000000 45211.000000 45211.000000
mean    40.936210 1362.272058   15.806419  258.163080   2.763841

```

std	10.618762	3044.765829	8.322476	257.527812	3.098021
min	18.000000	-8019.000000	1.000000	0.000000	1.000000
25%	33.000000	72.000000	8.000000	103.000000	1.000000
50%	39.000000	448.000000	16.000000	180.000000	2.000000
75%	48.000000	1428.000000	21.000000	319.000000	3.000000
max	95.000000	102127.000000	31.000000	4918.000000	63.000000

	pdays	previous
count	45211.000000	45211.000000
mean	40.197828	0.580323
	std	
	100.128746	2.303441
	min	
	-1.000000	0.000000
	25%	
	-1.000000	0.000000
	50%	
	-1.000000	0.000000
	75%	
	-1.000000	0.000000
	max	
	871.000000	275.000000

```
[9]: df.describe(include='all')
```

```
[9]:
```

	age	job	marital	education	default	balance \
count	45211.000000	45211	45211	45211	45211	45211.000000
unique	NaN	12	3	4	2	NaN
top	NaN	blue-collar	married	secondary	no	NaN
freq	NaN	9732	27214	23202	44396	NaN
mean	40.936210	NaN	NaN	NaN	NaN	1362.272058
std	10.618762	NaN	NaN	NaN	NaN	3044.765829
min	18.000000	NaN	NaN	NaN	NaN	-8019.000000
25%	33.000000	NaN	NaN	NaN	NaN	72.000000
50%	39.000000	NaN	NaN	NaN	NaN	448.000000
75%	48.000000	NaN	NaN	NaN	NaN	1428.000000
max	95.000000	NaN	NaN	NaN	NaN	102127.000000

housing loan

contact day month duration \

count	45211	45211	45211	45211.000000	45211	45211.000000
unique	2	2	3	NaN	12	NaN
top	yes	no	cellular	NaN	may	NaN
freq	25130	37967	29285	NaN	13766	NaN
mean	NaN	NaN	NaN	15.806419	NaN	258.163080
std	NaN	NaN	NaN	8.322476	NaN	257.527812
min	NaN	NaN	NaN	1.000000	NaN	0.000000
25%	NaN	NaN	NaN	8.000000	NaN	103.000000
50%	NaN	NaN	NaN	16.000000	NaN	180.000000
75%	NaN	NaN	NaN	21.000000	NaN	319.000000
max	NaN	NaN	NaN	31.000000	NaN	4918.000000

campaign pdays

previous poutcome y

count	45211.000000	45211.000000	45211.000000	45211	45211
unique	NaN	NaN	NaN	4	2

top	NaN	NaN	NaN	unknown	no
freq	NaN	NaN	NaN	36959	39922
mean	2.763841	40.197828	0.580323	NaN	NaN
std	3.098021	100.128746	2.303441	NaN	NaN
min	1.000000	-1.000000	0.000000	NaN	NaN
25%	1.000000	-1.000000	0.000000	NaN	NaN
50%	2.000000	-1.000000	0.000000	NaN	NaN
75%	3.000000	-1.000000	0.000000	NaN	NaN
max	63.000000	871.000000	275.000000	NaN	NaN

```
[10]: count_duplicated = df.duplicated().sum()
print('Dataset having',count_duplicated,'duplicated
values')
```

Dataset having 0 duplicated values

```
[11]: cat_var=[]
for var in df.columns:
    if df[var].dtype=='object':
        cat_var.append(var)
categorical_variables=np.array(cat_var)
categorical_variables
```

```
[11]: array(['job', 'marital', 'education', 'default', 'housing', 'loan',
            'contact', 'month', 'poutcome', 'y'], dtype='<U9')
```

```
[12]: num_var=[]
for var in df.columns:
    if df[var].dtype=='int64':
        num_var.append(var)
numerical_variables=np.array(num_var)
numerical_variables
```

```
[12]: array(['age', 'balance', 'day', 'duration', 'campaign', 'pdays',
            'previous'], dtype='<U8')
```

```
[13]: for var in df.columns:
        print(df[var].value_counts())
```

```
32    2085
31    1996
33    1972
34    1930
35    1894
```

```
...
93     2
90     2
95     2
88     2
94     1
```

Name: age, Length: 77, dtype: int64

```

blue-collar    9732
management    9458
technician     7597
admin.         5171
services       4154
retired        2264
self-employed  1579
entrepreneur   1487
unemployed     1303
housemaid      1240
student        938
unknown        288
Name: job, dtype:
int64 married
  27214 single
  12790 divorced
   5207
Name: marital, dtype:
int64 secondary  23202
tertiary        13301
primary         6851
unknown         1857
Name: education, dtype: int64
no 44396 yes
   815
Name: default, dtype: int64
  0      3514
  1      195
  2      156
  4      139
  3      134
  ...
-381      1
  4617      1
 20584      1
  4358      1
 16353      1
Name: balance, Length: 7168, dtype:
int64 yes  25130 no   20081
Name: housing, dtype: int64
no   37967
yes   7244
Name: loan, dtype:
int64 cellular
  29285 unknown
  13020 telephone  2906

```

Name: contact, dtype: int64

20	2752
18	2308
21	2026
17	1939
6	1932
5	1910
14	1848
8	1842
28	1830
7	1817
19	1757
29	1745
15	1703
12	1603
13	1585
30	1566
9	1561
11	1479
4	1445
16	1415
2	1293
27	1121
3	1079
26	1035
23	939
22	905
25	840
31	643
10	524
24	447
1	322

Name: day, dtype:

int64	may	13766
jul	6895	aug
6247	jun	5341 nov
3970	apr	2932 feb
2649	jan	1403 oct
738	sep	579 mar
477	dec	214

Name: month, dtype: int64

124	188
90	184
89	177
104	175
122	175

```

...
1833      1
1545      1
1352      1
1342      1
1556      1
Name: duration, Length: 1573, dtype: int64
1      17544
2      12505 3      5521
4      3522
5      1764
6      1291
7      735
8      540
9      327
10     266
11     201
12     155
13     133
14     93
15     84
16     79
17     69
18     51
19     44
20     43
21     35
22     23
25     22
23     22
24     20
29     16
28     16
26     13
31     12
27     10
32     9
30     8
33     6
34     5
36     4
35     4
43     3
38     3
37     2
50     2

```

41	2
46	1
58	1
55	1
63	1
51	1
39	1
44	1

Name: campaign, dtype: int64

-1	36954
182	167
92	147
91	126
183	126
...	
449	1
452	1
648	1
595	1
530	1

Name: pdays, Length: 559, dtype: int64

0	36954
1	2772
2	2106
3	1142
4	714
5	459
6	277
7	205
8	129
9	92
10	67
11	65
12	44
13	38
15	20
14	19
17	15
16	13
19	11
20	8
23	8
18	6
22	6
24	5
27	5

21	4
29	4
25	4
30	3
38	2
37	2
26	2
28	2
51	1
275	1
58	1
32	1
40	1
55	1
35	1
41	1

Name: previous, dtype: int64


unknown	36959
failure	4901
other	1840
success	1511

Name: poutcome, dtype: int64


no	39922
yes	5289

Name: y, dtype: int64

```
[14]: data={
    'blue-collar':9732,
    'management':9458,
    'technician':7597,
    'admin':5171,
    'services':4154,
    'retired':2264,
    'self-employed':1579,
    'entrepreneur':1487,
    'unemployed':1303,
    'housemaid':1240,
    'student':938,
    'unknown':288
}
x_labels=list(data.keys())
y_labels=list(data.values())
plt.figure(figsize=(10, 6))
plt.barh(x_labels, y_labels, color='skyblue')
plt.xlabel('Count')
plt.ylabel('Occupation')
plt.title('Count of Each Occupation')
plt.show()
```




```
[15]: dict={
    'married':27214,
    'single':12790,
    'divorced':5207
}
label_x=list(dict.keys())
label_y=list(dict.values())
plt.figure(figsize=(10, 6))
plt.barh(label_x,label_y, color='skyblue')
plt.xlabel('Count')
plt.ylabel('Marital status')
plt.title('Description about their Maritial status ')
plt.show()
```




```
[16]: d={
    'secondary':23202,
    'tertiary':13301,
    'primary':6851,
    'unknown':1857
}
u=list(d.keys())
v=list(d.values())
plt.scatter(u,v,color='blue')
```

```
plt.plot(u,v,color='red')  
plt.show()
```



```
[17]: plt.barh(u,v,color='red')  
plt.xlabel('Count')  
plt.ylabel('Education Standard')  
plt.title('Education')  
plt.show()
```



```
[18]: df=df.replace('unknown',np.nan)
```

```
[19]: da=pd.DataFrame({'columns':df.columns,'number_of_nulls_values':df.isna().
    .sum(),'percentage_null_values':round(df.isna().sum()*100/len(df),2)})
```

```
[20]:
```


da

```
[20]:
```

	columns	number_of_nulls_values	percentage_null_values
age	age	0	0.00
job	job	288	0.64
marital	marital	0	0.00
education	education	1857	4.11
default	default	0	0.00
balance	balance	0	0.00
housing	housing	0	0.00
loan	loan	0	0.00
contact	contact	13020	28.80
day	day	0	0.00
month	month	0	0.00
duration	duration	0	0.00
campaign	campaign	0	0.00

pdays	pdays	0	0.00
previous	previous	0	0.00
poutcome	poutcome	36959	81.75
y	y	0	0.00

```
[21]: plt.figure(figsize=(10,7))
plt.barh('columns','percentage_null_values',data=da)
plt.xlabel('Percentage')
plt.ylabel('Columns')
plt.title('Plot for percentage of null values after replacing the unknowns ')
plt.show()
```



```
[22]: null_variables=['poutcome','contact','education','
job'] for x in null_variables:
    print(df[x].value_counts())
print('-----')
```

```

failure    4901
other      1840
success    1511
Name: poutcome, dtype: int64
-----
-----
---
```

```

cellular    29285
telephone   2906
Name: contact, dtype: int64
-----
-----
---
```

```

secondary   23202
tertiary    13301
primary     6851
Name: education, dtype: int64
-----
-----
---
```

```

blue-collar    9732
management     9458
technician     7597
admin.         5171
services       4154
retired        2264
self-employed  1579
entrepreneur   1487
unemployed     1303
housemaid      1240
student        938
Name: job, dtype: int64
-----
-----
---
```

```
[23]: df.drop(columns='poutcome', inplace=True)
      df.shape
```

```
[23]: (45211, 16)
```

```
[24]: df['contact']=df['contact'].fillna(df['contact'].mode()[0])
      df['education']=df['education'].fillna(df['education'].mode()[0])
      df['job']=df['job'].fillna(df['job'].mode()[0])
```

```
[25]: df.isna().sum()
```

```
[25]: age      0
```

```

job            0
marital        0
education      0
default        0
balance        0
housing        0
loan           0
contact        0
day            0
month          0
duration       0
campaign       0
pdays        0
previous       0
y              0
dtype: int64

```

```
[26]: columns_mean=df.mean()
      columns_mean
```



C:\Users\PHALGUN\AppData\Local\Temp\ipykernel_26316\4192548252.py:1:
FutureWarning: The default value of numeric_only in DataFrame.mean is deprecated. In a future version, it will default to False. In addition, specifying 'numeric_only=None' is deprecated. Select only valid columns or specify the value of numeric_only to silence this warning.



```
columns_mean=df.mean()
```



```
[26]: age            40.936210
      balance       1362.272058
      day           15.806419
      duration      258.163080
      campaign       2.763841
      pdays         40.197828
      previous      0.580323
      dtype: float64
```

```
[27]: for column in df.select_dtypes(include='int64'):
      plt.figure(figsize=(4,3))
      plt.boxplot(df[column], vert=False)
      plt.title(f'Boxplot of {column}')
      plt.show()


      plt.hist(df[column], bins=20, density=True, alpha=0.6, color='b')
      plt.title(f'Histogram of {column}')
      plt.xlabel(column)
      plt.ylabel('Density')
      plt.grid(True)
      plt.show()
```












Boxplot of duration





Histogram of duration









```
[28]: df['marital'] =  
df['marital'].map({'single':0, 'married':1, 'divorced':2})  
df['education'] = df['education'].map({'secondary':0, 'tertiary':1,  
'primary':2}) df['default'] = df['default'].map({'yes':1, 'no':0})
```

```
df['housing'] = df['housing'].map({'yes':1,'no':0}) df['loan'] =
df['loan'].map({'yes':1,'no':0}) df['contact'] =
df['contact'].map({'cellular':1,'telephone':0}) df['y'] =
df['y'].map({'yes':1,'no':0})
```

[29]: df

```
[29]:      age      job marital education default balance housing loan \
0      58  management 1      1      0      2143 1      0
1      44  technician 0      0      0      29  1      0 2    33
      entrepreneur 1      0      0      2      1      1
3      47  blue-collar 1      0      0      1506 1      0
4      33  blue-collar 0      0      0      1      0      0
...    ...
45206  51  technician 1      1      0      825 0      0
45207  71  retired   2      2      0      1729 0      0
45208  72  retired   1      0      0      5715 0      0
45209  57  blue-collar 1      0      0      668 0      0
45210  37  entrepreneur 1      0      0      2971 0      0

      contact day month duration campaign pdays previous y
0      1      5      may  261  1      -1      0 0 11      5
may  151  1      -1      0 0
2      1      5      may  76  1      -1      0 0
3      1 5 may 92 1 -1 0 0 4 1 5 may 198 1 -1 0 0 ... ..
... ..
45206      1 17 nov 977 3      -1      0 1
45207      1 17 nov 456 2      -1      0 1
45208      1 17 nov 1127 5      184  3 1
45209      0 17 nov 508 4      -1      0 0
45210      1 17 nov 361 2      188  11 0
```

[45211 rows x 16 columns]

```
[30]: df=pd.get_dummies(df, columns=['job', "month"], prefix=["job", "month"],
.drop_first=True)
```

[31]:

df

```
[31]:      age marital education default balance housing loan contact day \
0      58      1      1      0      2143 1      0      1      5 1 44      0      0
0      29      1      0      1      5 2 33      1      0      0      2      1      1
      1      5 3 47      1      0      0      1506 1      0      1      5
4 33 0 0 0 1 0 0 1 5 ... ..
45206  51      1      1      0      825 0      0      1      17
45207  71      2      2      0      1729 0      0      1      17
```

45208	72	1	0	0	5715	0	0	1	17
45209	57	1	0	0	668	0	0	0	17
45210	37	1	0	0	2971	0	0	1	17

	duration	...	month_dec	month_feb	month_jan	month_jul	month_jun	\
0	261	...	0	0	0	0	0	
1	151	...	0	0	0	0	0	
2	76	...	0	0	0	0	0	
3	92	...	0	0	0	0	0	
4	198	...	0	0	0	0	0	
...	
45206	977	...	0	0	0	0	0	
45207	456	...	0	0	0	0	0	
45208	1127	...	0	0	0	0	0	
45209	508	...	0	0	0	0	0	
45210	361	...	0	0	0	0	0	

	month_mar	month_may	month_nov	month_oct	month_sep
0	0	1	0	0	0
1	0	1	0	0	0
2	0	1	0	0	0
3	0	1	0	0	0
4	0	1	0	0	0
...
45206	0	0	1	0	0
45207	0	0	1	0	0
45208	0	0	1	0	0
45209	0	0	1	0	0
45210	0	0	1	0	0

[45211 rows x 35 columns]

```
[32]: df.shape
```

```
[32]: (45211, 35)
```

```
[33]: df.info()
```

```
<class
'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to
```


45210 Data columns (total 35 columns):

#	Column	Non-Null	Count	Dtype
0	age	45211	non-null	int64
1	marital	45211	non-null	int64
2	education	45211	non-null	int64
3	default	45211	non-null	int64
4	balance	45211	non-null	int64
5	housing	45211	non-null	int64
6	loan	45211	non-null	int64
7	contact	45211	non-null	int64
8	day	45211	non-null	int64
9	duration	45211	non-null	int64
10	campaign	45211	non-null	int64
11	pdays	45211	non-null	int64
12	previous	45211	non-null	int64
13	y	45211	non-null	int64
14	job_blue-collar	45211	non-null	uint8
15	job_entrepreneur	45211	non-null	uint8
16	job_housemaid	45211	non-null	uint8
17	job_management	45211	non-null	uint8
18	job_retired	45211	non-null	uint8
19	job_self-employed	45211	non-null	uint8
20	job_services	45211	non-null	uint8
21	job_student	45211	non-null	uint8

22	job_technician	45211	non-null
		uint8	
23	job_unemployed	45211	non-null
		uint8	
24	month_aug	45211	non-null
		uint8	
25	month_dec	45211	non-null
		uint8	
26	month_feb	45211	non-null
		uint8	
27	month_jan	45211	non-null
		uint8	
28	month_jul	45211	non-null
		uint8	
29	month_jun	45211	non-null
		uint8	
30	month_mar	45211	non-null
		uint8	
31	month_may	45211	non-null
		uint8	
32	month_nov	45211	non-null
		uint8	
33	month_oct	45211	non-null
		uint8	
34	month_sep	45211	non-null
		uint8	

dtypes: int64(14), uint8(21)

memory usage: 5.7 MB

```
[34]: dependent_variable='y' independent_variables =
list(set(df.columns.tolist()) - {dependent_variable})
X =
df[independent_variables].copy()
y =
df[dependent_variable].copy()
```

```
[35]: from sklearn.model_selection import train_test_split
```

```
[36]: x_train,x_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=101)
```

```
[37]: from sklearn.neighbors import KNeighborsClassifier
      from sklearn.svm import SVC
      from sklearn.naive_bayes import CategoricalNB
      from sklearn.cluster import KMeans
      from sklearn.model_selection import cross_val_predict
      from tabulate import tabulate
```

```
[38]: knn = KNeighborsClassifier()
```

```
[39]: knn.fit(x_train, y_train)
```

```
[39]: KNeighborsClassifier()
```

```
[40]: knn_pred = knn.predict(x_test)
```

```
[41]: from sklearn.metrics import accuracy_score, precision_score, recall_score, _
      f1_score
```

```
[42]: from sklearn.metrics import confusion_matrix,classification_report
```

```
[43]: print(confusion_matrix(y_test,knn_pred))
```

```
[[7668  276]
 [ 803 296]]
```

```
[44]: knn_report=classification_report(y_test,knn_pred)
      print(knn_report)
```

	precision	recall	f1-score	support
0	0.91	0.97	0.93	7944
1	0.52	0.27	0.35	1099
accuracy			0.88	9043
macro avg	0.71	0.62	0.64	9043
weighted avg	0.86	0.88	0.86	9043

```
[45]: svm = SVC()
```

```
[46]: svm.fit(x_train,y_train)
```

```
[46]: SVC()
```

```
[47]: svm_pred = svm.predict(x_test)
```

```
[48]: print(confusion_matrix(y_test,svm_pred))
```

```
[[7941  3]
 [1097  2]]
```

```
[49]: svm_report = classification_report(y_test,svm_pred)
print(svm_report)
```

```

              precision    recall  f1-score   support

    0           0.88         1.00         0.94         7944
    1           0.40         0.00         0.00         1099

 accuracy               0.88         9043
 macro avg           0.64         0.50         0.47         9043
 weighted avg        0.82         0.88         0.82         9043
```

```
[50]: from sklearn.linear_model import LogisticRegression
```

```
[51]: model = LogisticRegression()
```

```
[52]: model.fit(x_train,y_train)
```

```
C:\Users\PHALGUN\anaconda3\lib\sitepackages\sklearn\linear_model\_log
istic.py:458: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in: <https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

```

              https://scikit-
              learn.org/stable/modules/linear_model.html#logistic-
regression
    n_iter_i = _check_optimize_result(
```

```
[52]: LogisticRegression()
```

```
[53]: pred = model.predict(x_test)
```

```
[54]: print(confusion_matrix(y_test,pred))
```

```
[[7828 116]
 [ 890 209]]
```

```
[55]: print(classification_report(y_test,pred))
```

	precision	recall	f1-score	support
0	0.90	0.99	0.94	7944
1	0.64	0.19	0.29	1099
accuracy			0.89	9043
macro avg	0.77	0.59	0.62	9043
weighted avg	0.87	0.89	0.86	9043

```
[56]: kmeans = KMeans(n_clusters=2)
```

```
[57]: kmeans.fit(x_train)
```

```
C:\Users\PHALGUN\anaconda3\lib\site-  
packages\sklearn\cluster\_kmeans.py:870:  
FutureWarning: The default value of `n_init` will change from 10 to  
'auto' in  
1.4. Set the value of `n_init` explicitly to suppress the warning  
warnings.warn(
```

```
[57]: KMeans(n_clusters=2)
```

```
[58]: train_cluster_labels = kmeans.predict(x_train)  
test_cluster_labels = kmeans.predict(x_test)
```

```
[59]: def map_cluster_labels(cluster_labels,  
    true_labels): cluster_to_label = {}  
    for cluster in set(cluster_labels): idx =  
        cluster_labels == cluster  
        label = true_labels[idx].mode()[0]  
        cluster_to_label[cluster] = label  
    mapped_labels = [cluster_to_label[cluster] for cluster in  
        cluster_labels]  
    return mapped_labels
```

```
[60]: y_train_kmeans_mapped = map_cluster_labels(train_cluster_labels,  
    y_train)  
y_test_kmeans_mapped = map_cluster_labels(test_cluster_labels, y_test)
```

```
[61]: print(confusion_matrix(y_test,y_test_kmeans_mapped))
```

```
[[7944  0]  
 [1099  0]]
```

```
[62]: kmc_report=classification_report(y_test, y_test_kmeans_mapped)  
print(kmc_report)
```

	precision	recall	f1-score	support
0	0.88	1.00	0.94	7944
1	0.00	0.00	0.00	1099
accuracy	0.88	9043	macro avg	0.44 0.50 0.47
9043	weighted avg	0.77	0.88	0.82 9043

C:\Users\PHALGUN\anaconda3\lib\sitepackages\sklearn\metrics_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

C:\Users\PHALGUN\anaconda3\lib\sitepackages\sklearn\metrics_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

C:\Users\PHALGUN\anaconda3\lib\sitepackages\sklearn\metrics_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

```
[63]: models={
    "SVM":svm,
    "KNN":knn,
    "LOGISTICREGRESSION":model,
    "KMeans":kmeans,
}
```

```
[64]: accuracy={}
precision={}
recall={}
f1={}
```

```
[65]: for name, model in models.items():
    y_pred = cross_val_predict(model, x_test, y_test, cv=5)
    accuracy[name] = accuracy_score(y_test, y_pred)
    precision[name] = precision_score(y_test, y_pred, average='weighted')
    recall[name] = recall_score(y_test, y_pred, average='weighted')
    f1[name] = f1_score(y_test, y_pred, average='weighted')
```

C:\Users\PHALGUN\anaconda3\lib\sitepackages\sklearn\linear_model_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

[learn.org/stable/modules/linear_model.html#logistic-](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
regression

```
n_iter_i = _check_optimize_result(  
C:\Users\PHALGUN\anaconda3\lib\sitepackages\sklearn\linear_model\_log  
istic.py:458: ConvergenceWarning: lbfgs failed to converge  
(status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in: [https://scikit-](https://scikit-learn.org/stable/modules/preprocessing.html)

[learn.org/stable/modules/preprocessing.html](https://scikit-learn.org/stable/modules/preprocessing.html)

Please also refer to the documentation for alternative solver options:

[https://scikit-](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

[learn.org/stable/modules/linear_model.html#logistic-](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
regression

```
n_iter_i = _check_optimize_result(  
C:\Users\PHALGUN\anaconda3\lib\sitepackages\sklearn\linear_model\_log  
istic.py:458: ConvergenceWarning: lbfgs failed to converge  
(status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in: [https://scikit-](https://scikit-learn.org/stable/modules/preprocessing.html)

[learn.org/stable/modules/preprocessing.html](https://scikit-learn.org/stable/modules/preprocessing.html)

Please also refer to the documentation for alternative solver options:

[https://scikit-](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

[learn.org/stable/modules/linear_model.html#logistic-](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
regression

```
n_iter_i = _check_optimize_result(  
C:\Users\PHALGUN\anaconda3\lib\sitepackages\sklearn\linear_model\_log  
istic.py:458: ConvergenceWarning: lbfgs failed to converge  
(status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in: [https://scikit-](https://scikit-learn.org/stable/modules/preprocessing.html)

[learn.org/stable/modules/preprocessing.html](https://scikit-learn.org/stable/modules/preprocessing.html)

Please also refer to the documentation for alternative solver options:

```
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
C:\Users\PHALGUN\anaconda3\lib\sitepackages\sklearn\linear_model\_logistic.py:458: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in: <https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

```
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
C:\Users\PHALGUN\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
C:\Users\PHALGUN\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
C:\Users\PHALGUN\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
C:\Users\PHALGUN\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
C:\Users\PHALGUN\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:870:
```


FutureWarning: The default value of `n_init` will change from 10 to 'auto' in

1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(

```
[66]: fig, axes = plt.subplots(2, 2, figsize=(7,5))


axes[0, 0].bar(accuracy.keys(), accuracy.values())
axes[0, 0].set_title('Accuracy')
axes[0, 1].bar(precision.keys(), precision.values())
axes[0, 1].set_title('Precision')
axes[1, 0].bar(recall.keys(), recall.values())
axes[1, 0].set_title('Recall')
axes[1, 1].bar(f1.keys(), f1.values())
axes[1, 1].set_title('F1 Score')

for ax in axes.flat:
    ax.set_xticklabels(ax.get_xticklabels(), rotation=45, ha='right')


plt.tight_layout()
plt.show()
```



C:\Users\PHALGUN\AppData\Local\Temp\ipykernel_26316\785931337.py:13:


UserWarning: FixedFormatter should only be used together with
FixedLocator ax.set_xticklabels(ax.get_xticklabels(), rotation=45,
ha='right')



```
[67]: for name in models.keys(): fig, axes
      = plt.subplots(figsize=(6,4))
          axes.bar(["Accuracy", "Precision", "Recall", "F1 Score"],
                  [accuracy[name],
                    precision[name], recall[name],
                    f1[name]]) axes.set_title(name)
      plt.show()
```







```
[68]: metrics_table = []
      for name in models.keys():
          metrics_table.append({
              "Model": name,
              "Accuracy": accuracy[name],
              "Precision": precision[name],
              "Recall": recall[name],
              "F1 Score": f1[name]
          })
```

```
[69]: print(tabulate(metrics_table, headers="keys", tablefmt="grid"))
```

```
+-----+-----+-----+-----+-----+
---+
| Model          | Accuracy | Precision | Recall | F1 Score |
+=====+=====+=====+=====+=====+
===+
| SVM            | 0.879575 | 0.850356 | 0.879575 | 0.827157 |
+-----+-----+-----+-----+-----+
---+
| KNN            | 0.881455 | 0.859081 | 0.881455 | 0.864535 |
+-----+-----+-----+-----+-----+
---+
| LOGISTICREGRESSION | 0.890412 | 0.869769 | 0.890412 | 0.867771 |
+-----+-----+-----+-----+-----+
---+
| KMeans         | 0.842088 | 0.793948 | 0.842088 | 0.814806 |
```

+-----+-----+-----+-----+-----
---+

[]:

[]:

[]:

[]: